


Predicting March Madness Basketball Games

Jonathen Marin

California State University, Fresno

May 6, 2026



Outline

- 1 Background
- 2 Data
- 3 Past Strategies
- 4 Modeling Techniques
- 5 Methods
- 6 2025 Results
- 7 2026 Results
- 8 References

Section 1

Background

Introduction

- March Madness is the NCAA Division I Basketball Tournament, a single-elimination bracket of 68 teams held each March and April across both the Men's and Women's divisions to crown the National Champions.
- The regular season determines the seeding. Teams are seeded 1 through 16 within each of four regions, where lower seeds are considered the strongest and are rewarded with matchups against the highest seeds.
- With 67 games decided by single elimination, upsets are frequent, which is why the tournament has become famous for its unpredictability.

Tournament Bracket




Figure 1: 2026 Bracket Example

- Kaggle is an online platform for data science and machine learning competitions, owned by Google.
- Companies, researchers, and organizations post prediction problems with datasets, evaluation metrics, and prize pools.
- Participants submit models that are scored on a hidden test set, with a public leaderboard ranking submissions in real time.
- Kaggle has become a benchmark in the ML community for evaluating predictive modeling approaches against a global field of competitors.

March Machine Learning Mania

- Annual Kaggle competition running since 2014, challenging participants to predict outcomes of every possible matchup in the NCAA tournament.
- Submissions are a probability $p \in [0, 1]$ for each of the 4556 possible matchups across Men's and Women's.
- Scored using the Brier Score:

$$BS = \frac{1}{N} \sum_{t=1}^N (f_t - o_t)^2 \quad (1)$$

- Where f_t is our prediction and o_t is the outcome.
- o_t is 1 if the team wins, and 0 if they lose.

Brier Score

- A naive 'coin flip' strategy (50% probability for all games) yields a Brier Score of 0.25, serving as our baseline.

Predicted	Outcome	Scenario	Brier Score
0.50	Win or Loss	Baseline	0.2500
1.00	Win	Perfectly confident, correct	0.0000
1.00	Loss	Perfectly confident, incorrect	1.0000
0.70	Win	Moderately confident, correct	0.0900
0.70	Loss	Moderately confident, incorrect	0.4900
0.90	Win	Highly confident, correct	0.0100
0.90	Loss	Highly confident, incorrect	0.8100

Table 1: Brier Score Examples

Prediction Difficulty

- Top predictive models plateau near 75% accuracy [1].
- Odds of a perfect bracket: 1 in 9.2 quintillion (random) to 1 in 128 billion (informed).
- No publicly verified bracket has survived past the Sweet Sixteen [2].
- \$3.1 billion wagered in 2025, sportsbook lines are the benchmark to beat [3].

Section 2

Data

Kaggle Data

- Kaggle provides 37 tables spanning over three decades, separated by gender.
- Contains detailed regular season and playoff results, plus tournament seeding tables.
- Detailed box scores begin in 2003 for men and 2010 for women.

Variable	Description
PTS	Points scored
FGM, FGA	Field goals made and attempted
FGM3, FGA3	Three-point field goals made and attempted
FTM, FTA	Free throws made and attempted
OR, DR	Offensive and defensive rebounds
Ast	Assists
TO	Turnovers
Blk	Blocks
PF	Personal fouls

Table 2: Detailed Result Variables

External Data Sources

- Box scores miss nuances such as pace of play and strength of schedule.
- Supplement Kaggle data with advanced efficiency metrics:
 - **KenPom** (Men's): industry standard, per-possession efficiency adjusted for opponent [4].
 - **Bart Torvik** (Women's): comparable methodology, used since KenPom is men's only [5].

Metric	Description
NetRtg	ORtg – DRtg
OffRtg	Pts per 100 poss.
DefRtg	Pts allowed per 100 poss.
AdjTempo	Possessions per 40 min

Table 3: KenPom Metrics

Metric	Description
ADJOE	Adjusted offensive efficiency
ADJDE	Adjusted defensive efficiency
AdjTempo	Adjusted tempo

Table 4: Bart Torvik Metrics

Section 3

Past Strategies

- Tournament seeding alone is a strong predictor via logistic regression with z-score transformation [6].
- Combined logistic regression and Markov chain models, benchmarked against nearly 100 ranking systems, established logistic regression as the baseline standard [7].
- No classifier (decision trees, random forests, naive Bayes, MLPs) exceeds 74–75% accuracy on NCAAB data [1].
- **Key takeaway:** feature construction matters more than model choice.

- **Fundamental Analysis:** construct a “true” probability independent of public sentiment.
 - Elo ratings [8], KenPom, FiveThirtyEight, Massey
 - Treats expert outputs as features, not final predictions
- **Efficient Market Hypothesis** [9]: betting markets aggregate information better than any single model.
 - Uses closing lines from major bookmakers.
 - `goto_conversion` algorithm [10] corrects for favorite-long shot bias.

- **Boosting** [11]: tiered features (seeding to efficiency to Elo to GLM team quality), implemented in XGBoost. Multiple top 10 finishes in 2024 and 2025 cite his work as the baseline for their submission.
- **Simulation pipelines** among top 2024 finishers: predict point margins, then simulate brackets thousands of times via normal distribution over margins.
- Inspired the three models in this paper: **Logistic Regression**, **XGBoost**, and **Linear Regression with Monte Carlo**.

Section 4

Modeling Techniques

Logistic Regression

- Binary classification model that applies the sigmoid function to a linear combination of features.

$$P(Y = 1 \mid \mathbf{x}) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x_1 + \dots + \beta_k x_k)}} \quad (2)$$

- $\mathbf{x} = (x_1, \dots, x_k)$: feature vector
- $\beta = (\beta_0, \dots, \beta_k)$: coefficients estimated via maximum likelihood
- Output is bounded between $[0, 1]$, making it interpretable as a probability.
- Assumes linear relationship between features and log-odds; nonlinearity addressed via polynomial or interaction terms.

Linear Regression

- Models a continuous outcome as a linear combination of features, with coefficients estimated by ordinary least squares [1].

$$\hat{Y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \cdots + \beta_k x_k \quad (3)$$

- In sports prediction, used to estimate point margins which are then converted to win probabilities via:
 - logistic transformation of the predicted margin, or
 - Monte Carlo simulation over the margin distribution.

Visualization




Figure 2: Logistic Regression Illustration [2]




Figure 3: Linear Regression Illustration[3]

Monte Carlo Simulation

- Computational technique that uses repeated random sampling to estimate probabilities for systems with uncertainty.
- Generates N independent samples $x_1, x_2, \dots, x_N$ from a distribution $f(x)$, then estimates probabilities by counting how often a condition is met:

$$\hat{P} = \frac{1}{N} \sum_{i=1}^N \mathbf{1}(x_i \in A) \quad (4)$$

- $\mathbf{1}(\cdot)$: indicator function (1 if condition met, 0 otherwise).
- A : event of interest; $\hat{P}$: estimate of true probability $P(A)$.
- By the Law of Large Numbers [4], $\hat{P} \rightarrow P(A)$ as $N \rightarrow \infty$.

- **Extreme Gradient Boosting** [5] is an ensemble of decision trees built sequentially.
- Each new tree corrects the residual errors of the previous trees.
- Final prediction is the sum of contributions from all trees:

$$\hat{y}_i = \sum_{k=1}^K f_k(x_i) \quad (5)$$

- Widely used in competitive ML for structured data; captures complex interactions and nonlinearity automatically.

XGBoost Example




Figure 4: How XGBoost combines trees sequentially [6]

Cross Validation

- Evaluation technique that systematically rotates which portion of the data is held out, so every observation serves as both training and validation [7].
- In k -fold cross validation: partition data into k equally sized folds. Train k times, each time using $k - 1$ folds for training and the remaining fold for validation.
- Final score is the average performance across all k folds.

CV Example




Figure 5: Cross Validation vs. Regular Train/Test Split [8]

Section 5

Methods

Logistic Regression: Setup

- **Training boundary:** features computed from regular season data, model trained on tournament outcomes.
- To prevent data leakage and reflect roster turnover, the model uses a **rolling two-year window**.
- Example: to forecast the 2025 tournament, train only on 2023 and 2024 tournament games.
- **Motivation:** high roster turnover in college basketball means a program undergoes near-complete personnel transformation within 2–3 seasons.
- Older data does not reflect current team strength.

Variable Selection: The Pace Problem

- Raw box score totals are confounded by **pace**.
- High-tempo teams accumulate more counting stats without being more efficient.
- To compare teams fairly, statistics must be expressed as **rates** (per opportunity) rather than totals.
- Solution: Dean Oliver's **Four Factors** [9] — each metric normalized by its natural opportunity denominator.

The Four Factors

- Possession estimation from box score data:

$$\text{Poss} = (FGA - OR) + TO + 0.44 \cdot FTA \quad (6)$$

$$eFG\% = \frac{FGM + 0.5 \cdot FGM3}{FGA} \quad TOV\% = \frac{TO}{\text{Poss}} \quad (7)$$

$$ORB\% = \frac{OR}{OR + DR_{opp}} \quad FTR = \frac{FTM}{FGA} \quad (8)$$

- **Effective FG%:** adjusts for the added value of three-pointers.
- **Turnover %:** turnovers per possession.
- **Offensive Rebound %:** rebounding dominance per opportunity.
- **Free Throw Rate:** free throw frequency relative to FG attempts.

Final Feature Set

- Beyond the Four Factors, the model adds **Win Percentage** as a fifth feature.
- Win% aggregates effects the Four Factors miss directly:
 - Coaching decisions
 - Late-game execution
 - Opponent quality
- All five features computed as season averages from regular season data.
- Transformed into **differentials** between competing teams, centering the analysis on the relative mismatch.

Model Implementation

- Logistic regression implemented separately for Men's and Women's divisions.
- Why not one combined model with a gender indicator?
 - Differences in pace, physicality, and three-point volume.
 - The Four Factors carry structurally different predictive weight across divisions.
 - Forcing shared parameters would mask these differences.
- Each model optimizes its coefficients to the unique statistical footprint of its division.

Model Specification

- Logit link function maps linear combination of differentials to a probability in $(0, 1)$:


$$\begin{aligned}\text{logit}(\hat{P}) = \ln\left(\frac{P}{1-P}\right) &= \beta_0 + \beta_1(\text{eFG_diff}) + \beta_2(\text{TOV_diff}) \\ &\quad + \beta_3(\text{ORB_diff}) + \beta_4(\text{FTR_diff}) \\ &\quad + \beta_5(\text{Win_diff})\end{aligned}\tag{9}$$

- P : probability of a Team 1 victory.

Multicollinearity Check

- All five predictors fall below $VIF = 3$ in both division so no problematic collinearity.

Collinearity
High collinearity (VIF) may inflate parameter uncertainty



Collinearity
High collinearity (VIF) may inflate parameter uncertainty




Figure 6: Men's (left) and Women's (right) VIF Plots for 2025

Feature Correlations

- Pairwise correlations are modest in both divisions; consistent with VIF results.




Figure 7: Men's (left) and Women's (right) Correlation Matrix for 2025

Coefficient Estimates

Men's Model				
Variable	Est.	SE	p	Sig.
Intercept	0.4352	0.2062	0.0348	*
eFG Diff	7.7952	6.4160	0.2244	
TOV Pct Diff	-20.5106	12.1028	0.0901	.
ORB Pct Diff	12.3634	4.1219	0.0027	**
FTR Diff	-1.4307	4.5405	0.7527	
Win Diff	1.3253	1.7033	0.4365	

Women's Model				
Variable	Est.	SE	p	Sig.
Intercept	-0.0888	0.2055	0.6657	
eFG Diff	17.4458	5.5413	0.0016	**
TOV Pct Diff	-19.8212	8.7655	0.0237	*
ORB Pct Diff	12.4308	3.5127	0.0004	***
FTR Diff	3.9742	5.1691	0.4420	
Win Diff	-1.8468	2.2764	0.4172	

- ORB% drives wins in both divisions (~ 12.4) — universal driver.
- eFG% significant in women's, not men's.

Why XGBoost?

- Logistic regression assumes a **linear** relationship between differentials.
- Reality is nonlinear and full of interactions:
 - Offensive rebounding only converts to wins past a shooting efficiency threshold.
 - Linear models cannot capture this without explicit feature engineering.
- XGBoost advantages:
 - No distributional assumptions on features.
 - Robust to scale differences, no normalization needed (e.g. points scored vs. steals).
 - Built-in regularization handles modest sample sizes.

- **Symmetric duplication:** every game contributes two rows.
 - Row 1: winner as Team 1 (outcome = 1).
 - Row 2: loser as Team 1 (outcome = 0).
- **Three benefits:**
 - Model learns from both perspectives.
 - No positional bias toward Team 1 slot.
 - Exactly 50/50 class balance.
- **Pooled across divisions** (2003–2024): 4,552 rows, 139 features.
- Binary `men_women` indicator retained to capture structural differences.

Feature Set: Three Tiers

- **Tier 1: Seeding** - T_1_seed , T_2_seed , $Seed_diff$.
 - Including absolute seeds captures nonlinearity (1 vs. 2 $\neq$ 8 vs. 9).
- **Tier 2: Simple season statistics** - win%, points scored / allowed, point differential, plus pairwise differentials.
- **Tier 3: Strength of schedule adjusted box scores.**
 - Off_ and Def_ prefixes for team's own offense/defense.
 - SOS_ prefix for opponent quality faced.
 - All adjusted for overtime: $adjOT = \frac{40+5 \cdot NumOT}{40}$.
- Plus: **Elo ratings** (covered next).

Seeding as Predictive Signal




Figure 8: Point Differential Declines with Seed in Both Divisions.




Figure 9: larger seed gaps $\rightarrow$ larger margins of victory.

Custom Elo Rating System

- Captures team strength — updates after every game.
- Adapted from chess Elo [8], structure inspired by [10].
- Expected win probability for team A :

$$E_A = \frac{1}{1 + 10^{(R_B - R_A)/400}} \quad (10)$$

- R_A, R_B : current Elo ratings, 400: scaling constant.
- Rating update for team A after each game:

$$R_A = R_A + k(S_A - E_A) \quad (11)$$


- S_A : actual outcome (1 if win, 0 if loss), $k = 64$, maximum elo points exchanged.

Elo Carryover Across Seasons

- Rather than resetting each season, ratings carry forward partially:

$$R_{\text{new}} = 0.5 \cdot R_{\text{final}} + 750 \quad (12)$$


- 50% of historical rating preserved; 50% regresses toward 1500.
- Captures program quality (Duke, UConn) without letting ratings drift indefinitely.
- Final feature set includes T_1_Elo , T_2_Elo , and Elo_diff .



Cross Validation: Selecting the Number of Rounds

- Number of boosting rounds chosen via 5-fold cross validation on the full tournament dataset.
- For each fold:
 - Train on 80% of games, validate on the held-out 20%.
 - Track validation loss (binary cross-entropy) as trees are added.
- **Early stopping:** halt when validation loss does not improve for 50 consecutive rounds.
- Selection criterion: average validation loss across the 5 folds.
- **Optimal number of rounds:** 415.

Feature Importance



Linear Regression Monte Carlo

- Logistic regression and XGBoost predict win probability **directly** from features.
- LRMC takes a different approach: predict the **score**, then derive win probability from simulated outcomes.
- **Motivation:** predicting a final score is a more tractable regression problem than directly estimating a win probability.
- Residual standard deviation $\hat{\sigma}$ from the regression captures inherent unpredictability and drives the simulation stage.

Two-Stage Pipeline

- **Stage 1: Linear Regression (Score Prediction)**

- Predicts each team's expected points given offensive efficiency, opponent's defensive efficiency, and pace.
- σ quantifies score uncertainty.

- **Stage 2: Monte Carlo Simulation**

- Draw thousands of samples from each team's score distribution.
 - Win probability = proportion of simulations in which the team outscores its opponent.
- Naturally produces calibrated probabilities reflecting underlying score uncertainty.

Feature Set

- Built around **opponent-adjusted efficiency ratings** from external sources:
 - **Men's:** KenPom [4] — OffRating, DefRating, AdjTempo.
 - **Women's:** Bart Torvik [5] — AdjOE, AdjDE, AdjTempo.
- **Pairing structure:** each team's offensive profile is matched against its opponent's defensive profile.
 - Mirrors how scoring actually occurs — output depends on the matchup, not just team strength.
- Additional features:
 - Elo differential (custom implementation [10]).
 - Net rating differential (men's only — limitation of the women's model).

LRMC: Model Specification

- Linear regression fit separately for Men's and Women's via OLS.
- Training data: 10,326 rows (Men's), 10,828 rows (Women's).

Men's model:

$$\begin{aligned} \text{Points}_A = & \beta_0 + \beta_1 \text{OffRating}_A + \beta_2 \text{DefRating}_B + \beta_3 \text{AdjTempo}_A \\ & + \beta_4 \text{AdjTempo}_B + \beta_5 \text{Diff_NetRtg} + \beta_6 \text{Diff_Elo} + \varepsilon \end{aligned} \quad (13)$$

Women's model:

$$\begin{aligned} \text{Points}_A = & \beta_0 + \beta_1 \text{AdjOE}_A + \beta_2 \text{AdjDE}_B + \beta_3 \text{AdjTempo}_A \\ & + \beta_4 \text{AdjTempo}_B + \beta_5 \text{Diff_Elo} + \varepsilon \end{aligned} \quad (14)$$

- $\varepsilon \sim \mathcal{N}(0, \hat{\sigma}^2)$ — residual error estimated from training data.

Monte Carlo Simulation


- Fitted regression produces expected scores $\hat{\mu}_A$ and $\hat{\mu}_B$.
- Draw simulated scores from normal distributions:

$$\text{Points}_A \sim \mathcal{N}(\hat{\mu}_A, \hat{\sigma}^2), \quad \text{Points}_B \sim \mathcal{N}(\hat{\mu}_B, \hat{\sigma}^2) \quad (15)$$

- Win probability = proportion of simulations where Team A outscores Team B:


$$\hat{P}(\text{Team A wins}) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}[\text{Points}_{A,i} > \text{Points}_{B,i}] \quad (16)$$

LPMC Example



LRMC Example

Duke vs Houston Simulations



LPMC: Regression Assumptions

- Three core assumptions of multiple linear regression:
 - **Linearity**: relationship between features and target is linear.
 - **Normality**: residuals are approximately normally distributed.
 - **Homoscedasticity**: residual variance is constant across predicted values.
- **Linearity & normality**: verified via residuals vs. fitted and Q-Q plots — assumptions reasonably satisfied in both models.
- **Homoscedasticity**: Breusch-Pagan significant, but driven by high-leverage outliers and large n . OLS estimates remain unbiased;
- **Multicollinearity**: men's model has $VIF \approx 7.5$ between Diff_NetRtg and Diff_Elo (correlation 0.89). Women's model is clean (all $VIF < 5$).

LRMC Coefficients

Men's Model				
Variable	Estimate	Std. Error	p-value	Sig.
Intercept	-207.4000	5.4360	<0.001	***
Off Rating	0.6519	0.0216	<0.001	***
Def Rating	0.6421	0.0239	<0.001	***
AdjT Team	1.0830	0.0378	<0.001	***
AdjT Opp	1.0090	0.0378	<0.001	***
Diff NetRtg	0.0540	0.0181	0.0029	**
Diff Elo	0.0013	0.0008	0.1229	

Women's Model				
Variable	Estimate	Std. Error	p-value	Sig.
Intercept	-202.9000	3.9730	<0.001	***
Off Rating	0.7309	0.0136	<0.001	***
Def Rating	0.7866	0.0189	<0.001	***
AdjT Team	0.9605	0.0333	<0.001	***
AdjT Opp	0.8542	0.0331	<0.001	***
Diff Elo	0.0024	0.0006	<0.001	***

- Women's: $R^2 = 0.524$, $\hat{\sigma} = 9.30$. Men's: $R^2 = 0.379$, $\hat{\sigma} = 9.51$.
- Men's **Diff_Elo** insignificant due to collinearity with Diff_NetRtg.

Section 6




2025 Results

2025 Model Comparison

Model	Men's	Women's	Combined	Brier	AUC
Logistic Regression	70.10%	77.60%	73.90%	0.18401	0.8098
XGBoost	77.60%	80.60%	79.10%	0.12761	0.9215
Monte Carlo	77.61%	82.09%	79.85%	0.12857	0.9125
Kaggle Winner	—	—	—	0.10411	—

- Monte Carlo achieves the highest combined accuracy (79.85%).
- XGBoost and Monte Carlo essentially tied on Brier Score.

2025 ROC Curves



Section 7

2026 Results


2026: Override Strategy

- Applied 32 deterministic probability overrides to highly predictable matchups before submission.
- **Motivation:** medal contention requires aggressive submission, not just calibrated probabilities.
 - A 0.95 probability still accumulates Brier penalty even when the favorite wins.
 - Setting probability to 1.0 accepts upset risk for zero penalty on expected outcomes.
- **Men's overrides:** 1 and 2 seeds in Round 1; 1 seeds in Round 2.
- **Women's overrides:** 1 and 2 seeds in Rounds 1 and 2; 1 seeds in Round 3; UConn in Round 4.

Override Outcome



- 30 of 32 overrides correct (zero Brier penalty).
- 2 overrides incorrect (maximum penalty of 1.0 each).
- **Net effect on Brier Score: +0.00440 (worse).**
 - Override submission: 0.13237
 - Base (no overrides): 0.12797
- Two incorrect overrides outweighed the cumulative benefit of 30 correct ones.

2026 Confusion Matrices



- Men's: 74.60%. Women's: 80.95%. Combined: 77.78%.

2026 ROC Curves



- Base AUC (0.9021) outperforms Override AUC (0.8871).

2026 Competition Results

- **Base Monte Carlo:** 0.12797 Brier $\rightarrow$ 387th place.
- **Override submission:** 0.13237 Brier $\rightarrow$ 822nd place.
- Medal cutoff (8th place): 0.11845 — base model fell short by only 0.00352.
- Competition winner: 0.10975.
- **Bracket simulation results:**
 - 26 of 32 Round 1, 12 of 16 Round 2, 5 of 8 Sweet 16.
 - 2 of 4 Elite 8, 1 of 2 Final 4.
 - Correctly predicted Michigan (Men's) and UCLA (Women's) as champions in that iteration.

Section 8

References

References (1/2)



Zimmermann, A., Moorthy, S., & Shi, Z. (2013). Predicting College Basketball Match Outcomes Using Machine Learning Techniques. *arXiv:1310.3607*.



Harris, J. (2026). How Hard Is It to Pick a Perfect Bracket? *NCAA.com*.



Cohen, D. (2025). Americans to Legally Wager Estimated \$3.1 Billion on March Madness. *American Gaming Association*.



Pomeroy, K. (2025). *Pomeroy College Basketball Ratings*. The Forecast Factory LLC.



Torvik, B. (2017). T-Rank FAQ. *Blogspot*.



Schwertman, N. C., Schenk, K. L., & Holbrook, B. C. (1996). More Probability Models for the NCAA Regional Basketball Tournaments. *The American Statistician*, 50(1), 34-38.



Bracketology Driven by Data. (2026). *Georgia Tech News Center*.



Glickman, M. E., & Doan, T. (2025). *The US Chess Rating System*.



Burns, W. (2024). Efficient-Market Hypothesis (EMH). *EBSCO Information Services*.



gotoConversion. (2025). *Goto_conversion* [GitHub Repository].

References (2/2)

 Lumivero. (2025). Ordinary Least Squares Regression (OLS). *XLSTAT*.

 Kanade, V. (2022). Everything You Need to Know about Logistic Regression. *Spiceworks*.

 Raschka, S. *LinearRegression: An Implementation of Ordinary Least-Squares Linear Regression*. mlxtend.

 Law of Large Numbers. (2022). *EBSCO Information Services*.

 Kavlakoglu, E., & Russi, E. (2024). XGBoost. *IBM*.

 Yao, X., Fu, X., & Zong, C. (2022). Short-Term Load Forecasting Method Based on Feature Preference Strategy and LightGBM-XGboost. *IEEE Access*.

 Rudin, C. *Cross Validation Duke Course Notes*.

 Ferrer, J. (2025). What Is Cross-Validation? A Plain English Guide with Diagrams. *KDNuggets*.

 Oliver, D. (2004). *Basketball on Paper*. Potomac Books.

 Haupts, L. (2025). Calculate Elo Ratings. *Kaggle Notebook*.

Thank You!